

Winter 2020

# CS 133 Lab 5

## FPGA w/ OpenCL: Convolutional Neural Network (CNN)

Due: 3/12/20 10:00PM

### Description

Your task is to accelerate the computation of convolutional neural network (CNN) using OpenCL with a FPGA. The details of the algorithm can be found in the lecture slides and the Lab 3 description.

### Preparation

#### FPGA development and Project Overview

FPGA development roughly involves 3 steps: high-level synthesis (HLS) run, bitstream generation, and on-board execution. In this class, we only target HLS step. You either can do it using OpenCL as the entry point to be compiled by Xilinx's Vitis tool, or trying a flow with Falcon Computing's Merlin Compiler (using C + program approach similar to OpenMP). If you chose to do both, one will be counted as an extra credit for you so you can earn up to 110 additional bonus points. Please note that in order to get the bonus point for an extra lab:

1. You should **answer the comparison questions** (check out the submission part of Merlin) in your report.
2. You should **have received some points for that part's performance score**.

### Create an AWS Instance

Please create an AWS instance in US East (N. Virginia) for HLS step. You may choose any server with more than 32GB of memory (list available at <https://github.com/aws/aws-fpga>, FPGA Developer AMI section). TA used c5.4xlarge instance.

Please use your AWS account (not AWS Educate classroom account) for this lab. You will probably need to request an increase of limit for the type of instance you want. To do this, go to the AWS console -> Limits -> "Request limit increase" next to the instance type you want.

## <Vitis Part>

For Vitis flow, when asked to choose AMI, please click AWS MarketPlace and search for FPGA Developer AMI, as recommended in <https://github.com/aws/aws-fpga/tree/master/Vitis> . The version should be 1.8.0.

When launching the instance, please click “storage” and increase the /dev/sdb disk size to about 15GB. FPGA projects may require a large disk space.

## Login to AWS

FPGA Developer AMI uses CentOS. **Thus, your login ID should be centos, not ubuntu.**

## Xilinx Forum

You may visit Xilinx’s AWS forum to see existing questions and answers: <https://forums.aws.amazon.com/forum.jspa?forumID=243>.

## Environment Setup

Xilinx has provided the Vitis environment setup files. Please run the following commands:

```
git clone https://github.com/aws/aws-fpga.git $AWS_FPGA_REPO_DIR
cd $AWS_FPGA_REPO_DIR
source vitis_setup.sh
```

The git clone only needs to be done once, but the setup file would need to be sourced every time you login. The first run of the source command will install missing programs, so it will take longer than the following runs.

## Hello\_world Example

Please follow the instructions in the following github page:

<https://github.com/aws/aws-fpga/blob/master/Vitis/README.md>

You only need to follow the steps up to “SW emulation”. The instructions are copied below:

```
cd $VITIS_DIR/examples/xilinx/hello_world
make clean
make check TARGETS=sw_emu DEVICES=$AWS_PLATFORM all
```

You may also choose to follow the instructions from our guest lecture on Monday:

<https://github.com/xupgit/awslabs>

### Tips

- To resume a session in case you lose your connection, you can run screen after login.
- You can recover your session with screen -DRR if you lost your ssh connection.

- You should **stop** your instance if you are going back and resume your work in a few hours or days. Your data will be preserved but you will be charged for the [EBS storage](#) for \$0.10 per GB per month (with default settings).
- The instance types used for FPGA development are very **expensive**. The cost table can be found in <https://aws.amazon.com/ec2/pricing/on-demand/>. Please pay careful attention to your spendings.
- You should **terminate** your instance if you are not going to back and resume your work in days or weeks. **Data on the instance will be lost.**
- You are recommended to use **private** repos provided by [GitHub](#). **Do not put your code in a public repo.**

## Run CNN in Software Emulation Mode

We have prepared the host code for you at [GitHub](#).

The files are mostly the same as Lab 3 and 4 except changes to the makefile that enables FPGA development in Vitis environment. Also, the kernel file has been renamed to `xilinx.cl`. Log in to your Vitis instance and run the following commands:

```
git clone https://github.com/UCLA-VAST/cs-133-20w -o upstream
cd cs-133-20w/lab5
make swsim
```

The provided code will load test data and verify your results against a ground truth. It should run with a large error and finish in a few seconds.

We will use this software emulation mode for correctness checking only. Since this is an emulated result, **the execution time provided after running the software emulation is incorrect**. The performance estimation we will use for grading can be obtained from the HLS step of Vitis.

### Tips

- To check in your code to a **private** GitHub repo, [create a repo](#) first.

```
git branch -m upstream
git checkout -b master
git add xilinx.cl
git commit -m "lab5: first version" # change commit message accordingly
# please replace the URL with your own URL
git remote add origin git@github.com:YourGitHubUserName/your-repo-name.git
git push -u origin master
```

- You are recommended to `git add` and `git commit` often so that you can keep track of the history and revert whenever necessary.
- If you move to a new instance, just `git clone` your repo.
- Run `make test` to re-compile and test your code.
- If `make test` fails, it means your code produces the wrong result.
- ***Make sure your code produces correct results!***

## HLS Step using Vitis

Vitis provides high-level synthesis (HLS) step that synthesizes OpenCL kernel code into a hardware description (in Verilog or VHDL language). Along with hardware description generation, HLS step also provides performance estimation which we will use for performance grading. Please run

```
make hls
```

and you will be able to see the result in the following file:

```
./_x/cnn.hw.xilinx_aws-vu9p-f1-04261818_dynamic_5_0/CnnKernel/CnnKernel/solution/syn/report/CnnKernel_csynth.rpt
```

### Warning

- Due to the limited time left for grading, we will put a timeout for Vitis and Merlin HLS run. Your HLS step should be **completed within 90 minutes on a c5.4xlarge** machine, or you will receive **no performance grade**.

In this file, please scroll down to Performance Estimates -> Latency -> Summary -> Latency max column, which provides the number of cycles.

Assuming a default clock cycle of 250MHz, the execution time can be estimated as: (clock cycles) / 250MHz. The performance in terms of GFLOPS would be :  $kNum * kNum * kImSize * kImSize * kKernel * kKernel * 2 * 250 / (\text{clock cycles} * 1e3)$ .

### Warning

- Please make sure that all of the loops in your kernel file have ***fixed loop bounds*** (either by using constant variables or macros). They should not have a variable value. If any of the loop bounds is a variable, Xilinx HLS tool will not be able to provide an accurate performance estimate. If your max and min latency is different, we will grade your performance **based on the max latency**. If your latency is '?', you will receive **no performance grade**.

In the same file, you can find FPGA resource (LUT/FF/DSP48E/BRAM\_18K) usage. You can find it at Utilization Estimates -> Summary -> Total. Note that single-precision floating point adder uses 2 DSP48Es, single-precision floating point multiplier uses 3 DSP48Es, and 1 BRAM\_18K is 18Kbit. You will probably observe that as you apply more optimizations, your resource usage will increase. Ideally, you should try to keep applying optimization until your kernel occupies about 80% of these resources. The remaining 20% is reserved for interface (DRAM/PCIE controller) and other uses.

### Warning

- Please make sure that resource utilization is **less than 80% for all FPGA resources (LUT/FF/DSP48E/BRAM\_18K)**. If any of the resources are over this limit, you will receive **no performance grade**.

Your task is to implement a fast, parallel version of CNN on FPGA. You can start with the sequential version provided in `cnn.cpp`. You should edit `xilinx.cl` for this task. Please keep in

mind that, you are NOT allowed to adjust the global / local work size parameters in `params.sh` (they are fixed to 1). **Instead, parallelism should be exploited by using OpenCL FPGA directives.** The list of OpenCL directives supported in Vitis can be found in Chapter 2 of the following document:

[https://www.xilinx.com/support/documentation/sw\\_manuals/xilinx2018\\_2/ug1253-sdx-pragma-reference.pdf](https://www.xilinx.com/support/documentation/sw_manuals/xilinx2018_2/ug1253-sdx-pragma-reference.pdf)

For now, Vitis supports the same OpenCL directives as SDAccel.

## Example Code Snippet

In `xilinx.cl` file, we have provided an example code that consists of high-level description of the task and some code snippets which achieve throughput of  $ll=1$  for a  $5 \times 5$  tree reduction for the CNN convolution kernel (the `p` and `q` reduction loop). The code has been provided just for demonstration purpose, and it is up to you to decide whether to use whole or part of the code provided. Also, it is your task to attempt to convert it into a working version, and TAs will NOT be able to provide further assistance of the process.

## Submission

You need to report Vitis-estimated performance results of your FPGA-based OpenCL implementation on Xilinx Ultrascale+ VU9P FPGA. Please express your performance in GFlops and the speedup compared with the sequential version in CPU (you can reuse the result from Lab 3 or Lab 4). In particular, you need to submit a brief report which summarizes:

- (20 points) Please explain the parallelization and optimization strategies you have applied. Include the directives (if any) or code segments you have added to achieve this. Evaluate the performance of each parallelization/optimization that you have incrementally applied and explain why it improves the performance. Also, explain how your strategy differs from your Lab 3 CPU and Lab 4 GPU parallelization/optimization strategy and why you chose to apply a different strategy.
- (5 points) Please report the FPGA resource (LUT/FF/DSP/BRAM) usage, in terms of resource count and percentage of total. Which resource has been used most, in terms of percentage?  
(*Optional*: Analyze your original code and check if the DSP/BRAM resource usage matches your expectation. You only need to account for the single-precision floating-point adders, multipliers, and size of arrays. Please attach related code segments to your report and show your work on how you computed the expected number. Provide reasonable discussion on possible reasons if they differ significantly.)
- *Optional*: The challenges you faced, and how you overcame them.

You will need to submit your optimized kernel codes. Please do not modify or submit the host code. Please submit to CCLE. Please verify the **correctness** of your code before submission.

Your final submission should be a tarball which contains and only contains the following files:

<Your UID>.tar.gz

└ <Your UID>

└─ `xilinx.cl` (required if Vitis flow is chosen, optional otherwise)

- └ CnnKernel.cpp (required if Merlin flow is chosen, optional otherwise)
- └ lab5-report.pdf (required)

File lab5-report.pdf must be in PDF format. You should make the tarball by copying your lab5-report.pdf to the lab5 directory and running

make tar UID=<Your UID>. The Merlin kernel, i.e. CnnKernel.cpp will be included by the tar command if there exists a merlin/src/CnnKernel.cpp . If you made the tarball in other ways, you ***MUST*** put it in the lab5 directory and check by running make check UID=<Your UID>. The check script will tell you whether the Merlin kernel is included properly. The make tar command and make check command should be invoked in the lab5 directory. **You will not receive any grade if your submission doesn't match this pattern.**

## Grading Policy

### Submission Format

**Your submission will only be graded if it complies with the requirement. In case of missing reports, missing codes, or compilation error, you will receive 0 for the corresponding category/categories.**

### Correctness (50%)

Please check the correctness of your implementation.

### Performance (25%+5%)

The performance point will be added only if you have the correct result, so please prioritize the correctness over performance. Your performance will be evaluated based on the ranges of throughput (GFlops). We will set five ranges after evaluating all submissions and assign the points as follows:

- Range A++, better than Range A+ performance: 25 points + 5 points (bonus)
- Range A+, better than Range A performance: 25 points + 3 points (bonus)
- Range A GFlops [100, 150]: 25 points
- Range B GFlops [50, 100): 20 points
- Range C GFlops [15, 20): 15 points
- Range D GFlops [1, 15): 10 points
- Speed up lower than range D: 0 points
- 

### Report (25%)

Points may be deducted if your report misses any of the sections described above.

## <Flow Using the Merlin Compiler>

### Create an AWS Instance

For Merlin flow, please click “AMIs” first, and click “Private images”. You should see an AMI named “CS133\_Merlin-Compiler”. Launch this AMI like a normal AWS instance.

### Login to AWS

Merlin AMI uses CentOS. Thus, your login ID should be centos, not ubuntu.

### Environment Setup

For Merlin flow, please follow instructions in Prerequisite and Compile on AWS with Merlin section of

<https://github.com/falconcomputing/merlin-compiler/blob/master/On-Cloud/AWS/README.md>. Specifically, you should run

```
git clone https://github.com/falconcomputing/merlin-compiler.git
export XILINX_SDX=/opt/Xilinx/SDx/2019.1.op2552052/
source $XILINX_SDX/settings64.sh
source /opt/xilinx/xrt/setup.sh
```

Again, the git clone needs to be done once, but the environment setup may need to be done every time you login.

### VectorAdd Example

Please follow the instructions in Compile on AWS with Merlin section of

<https://github.com/falconcomputing/merlin-compiler/blob/master/On-Cloud/AWS/COMPILE.md>,

You only need to follow the instructions up to “Run your code on CPU”. In addition, try running the `mcc_esimate` flow to obtain cycle estimation. The instructions are copied below:

```
cd merlin-compiler/Examples/vectoradd/build/cpu_multi_core
make run
cd ../xilinx_mo
make mcc_estimate
```

“make run” will check the functionality/correctness of your code on CPU. “make `mcc_estimate`” applies code transformation and runs HLS to estimate the performance of your design.

### Run CNN in Software Emulation Mode

We have added a Merlin subdirectory in your Lab 5 project folder. Log in to your Merlin instance and run the following commands:

```
git clone https://github.com/UCLA-VAST/cs-133-20w -o upstream
cd cs-133-20w/lab5/merlin
make cpu run
```

The kernel file `src/CnnKernel.cpp` already has a working code and should finish with no error.

## HLS Step

Merlin Compiler uses SDAccel's HLS flow to synthesize its C file into hardware description file. Merlin Compiler reads SDAccel's performance estimation file and generates its own performance estimate, which we will use for performance grading. Please run

```
make xilinx mcc_estimate
```

and you will be able to see the result in the following file:

```
./build/xilinx_mo/merlin.rpt
```

In this file, please scroll down to Source code hierarchy -> CnnKernel -> AC column, which provides the number of cycles.

Your task is to implement a fast, parallel version of CNN on FPGA. You can start with the baseline implementation provided in `src/CnnKernel.cpp`. The parallelism should be exploited by using Merlin directives. The list can be found in the Merlin manual provided in CCLE.

## Submission

You need to report Merlin-estimated performance results of your FPGA-based OpenCL implementation on Xilinx Ultrascale+ VU9P FPGA. Please express your performance in GFlops and the speedup compared with the sequential version in CPU (you can reuse the result from Lab 3 or Lab 4). In particular, you need to submit a brief report which summarizes:

- Please explain the parallelization and optimization strategies you have applied. Include the directives (if any) or code segments you have added to achieve this. Evaluate the performance of each parallelization/optimization that you have incrementally applied.
- Please report the FPGA resource (LUT/FF/DSP/BRAM) usage, in terms of resource count and percentage of total. Which resource has been used most, in terms of percentage? This information can be found in the SDAccel HLS report : `build/xilinx_mo/.merlin_prj/run/report/CnnKernel_csynth.rpt`
- Please analyze Merlin log file (`build/xilinx_mo/merlin.log`), Merlin report file (`build/xilinx_mo/merlin.rpt`), and SDAccel HLS report (`build/xilinx_mo/.merlin_prj/run/report/CnnKernel_csynth.rpt`). Try to find the optimization Merlin has performed on your code. Explain why you reached this conclusion. Explain how such optimizations would increase the performance of your kernel.
- If you have done both of the flows, in order to get the score for both, you should also include (otherwise, you only get the report score for one flow):
  - Please concentrate on making comparisons with the Vitis version. That is, if the same strategy was applied as your Vitis version, please mention this in your report and keep the description very brief. If the coding style (or directive) has



changed from Vitis version, you will need to show the difference. For the newly applied optimizations, please elaborate and explain why it improves the performance. If some optimizations were removed, please mention this in your report.

- Compare the performance of Vitis flow and Merlin flow. Which one is better? Why? Does the result meet your expectation? If not, please discuss possible reasons.

Please write your report in the same file as Vitis report - lab5-report.pdf.

You will need to submit your optimized kernel codes. Please do not modify or submit the host code. Please submit to CCLE. Please verify the correctness of your code before submission.

For submitting Merlin kernel file, please read the submission instructions in the Vitis section.

## Grading Policy

### Submission Format

**Your submission will only be graded if it complies with the requirement. In case of missing reports, missing codes, or compilation error, you will receive 0 for the corresponding category/categories.**

### Correctness (50%)

Please check the correctness of your implementation.

### Performance (25%+5%)

The performance point will be added only if you have the correct result, so please prioritize the correctness over performance. Your performance will be evaluated based on the ranges of throughput (GFlops). We will set five ranges after evaluating all submissions and assign the points as follows:

- Range A++, better than Range A+ performance: 25 points + 5 points (bonus)
- Range A+, better than Range A performance: 25 points + 3 points (bonus)
- Range A GFlops [150, 200]: 25 points
- Range B GFlops [80, 150): 20 points
- Range C GFlops [30, 80): 15 points
- Range D GFlops [1, 30): 10 points
- Speed up lower than range D: 0 points

### Report (25%)

Points may be deducted if your report misses any of the sections described above.

## Academic Integrity

All work is to be done individually, and any sources of help are to be explicitly cited. Any instance of academic dishonesty will be promptly reported to the Office of the Dean of Students. Academic dishonesty, includes, but is not limited to, cheating, fabrication, plagiarism, copying code from other students or from the internet or facilitating academic misconduct. We'll use automated software to identify similar sections between different student programming assignments or against popular Internet sources.

Students are not allowed to post the lab solutions on public websites (including Github).

Late policy: Late submission will be accepted for **12 hours** with a 10% penalty. No late submission will be accepted after that (i.e., work that is more than two days late will lose all the points).